

Self Managing Clusters at Scale

Kevin Downey | Eytan Avisror

Eytan Avisror

Staff Software Engineer, Intuit
Core contributor, Keikoproj
github @eytan-avisor



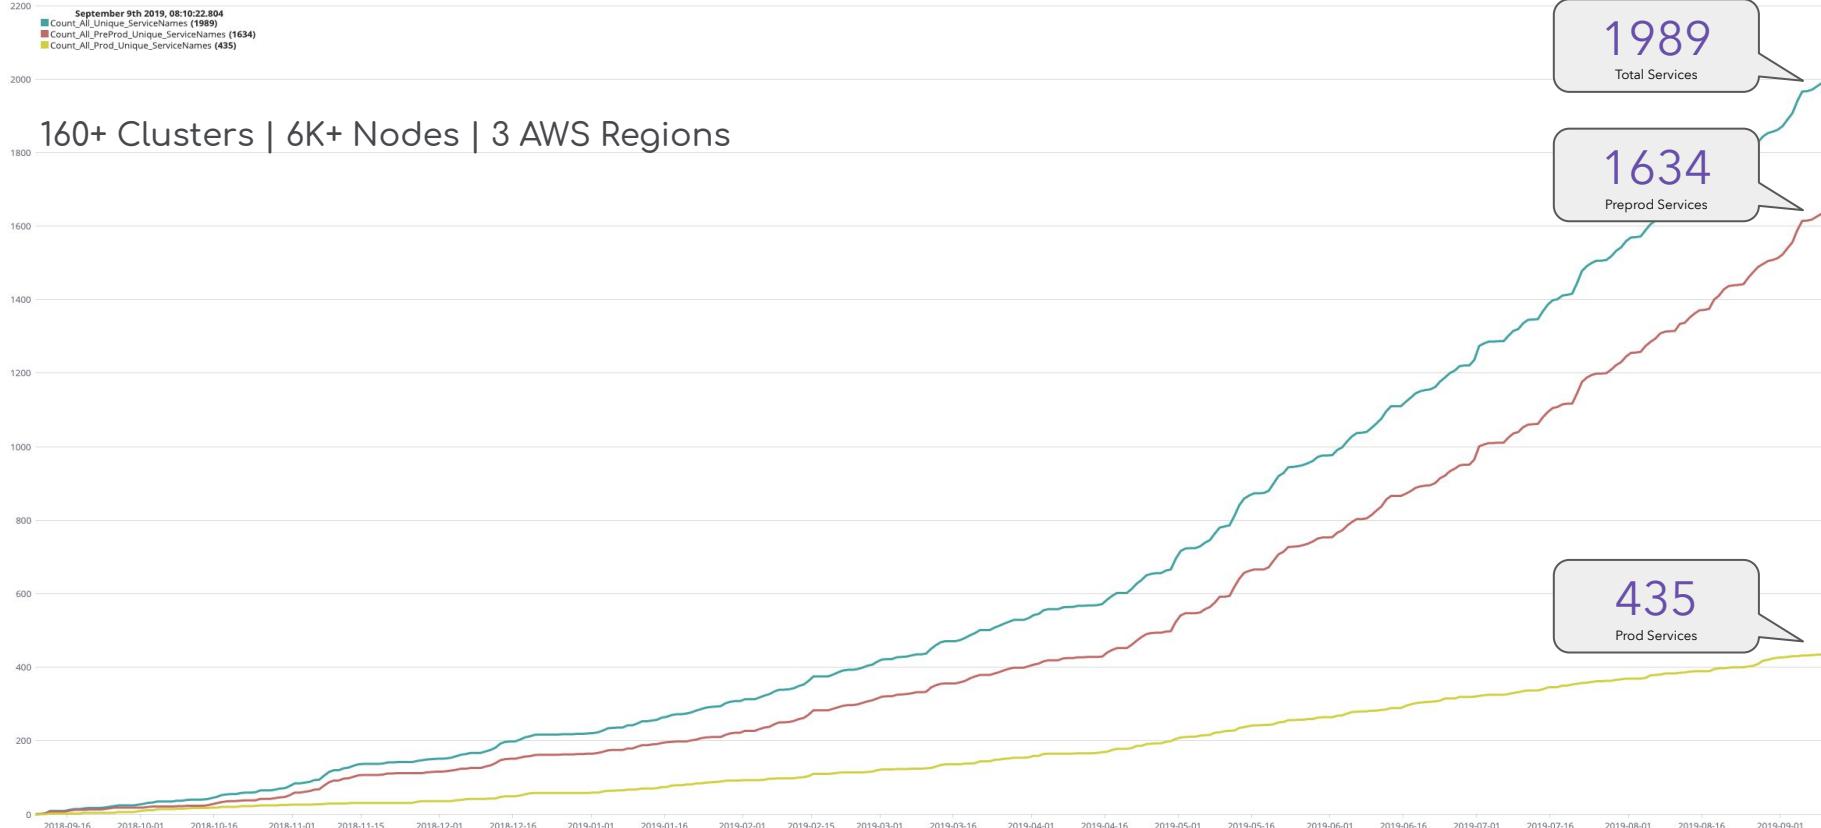
Kevin Downey

Staff Software Engineer, Intuit
Core contributor, Keikoproj
github @kevdowney



Kubernetes @ Intuit

Unique Services by PreProd/Prod



Modern Developer Platform

Kubernetes Powered, GitOps Driven

Kubernetes Powered, GitOps Driven



Application Repository

A git repository containing a set YAML and Kustomize files that make up your deployment.



Kubernetes Namespace

A share of a Kubernetes cluster using either a shared or dedicated node group.



Continuous Deployment with ArgoCD

ArgoCD enables GitOps between the repository and the Kubernetes cluster.



The Pipeline

Declarative build and deploy in seconds.

Common Problems

How do I manage nodes?

How do I manage critical cluster services?

How do I optimize the cost?

How do I ensure resiliency?

Keiko -
Enable Kubernetes at scale ...

What is Keiko?

Keiko provides a set of independent open-source tools for management of multi-tenant, reliable, secure and efficient Kubernetes clusters at scale.

github.com/keikoproj



Components

Orchestration

Addon-Manager Instance-Manager Upgrade-Manager

Security

Kube-Forensics

Reliability

Governor Lifecycle-Manager

Monitoring

Active-Monitor

Cost

Minion-Manager

Addon-Manager

How do I manage critical cluster addons?

What are cluster addons?

Critical Cluster Wide Services

1

Metrics and Logging

We want Metrics and Logging using Prometheus, Fluentd, Eventrouter, etc.

2

Networking

We want DNS and maybe mesh with CoreDNS, Istio, etc.

3

Autoscaling

We want to scale up cluster nodes using Cluster Autoscaler.

4

Ingress

We want to provide Ingress to our cluster using ALB Ingress Controller, Nginx Ingress Controller, etc.

Cluster Addon Management

Declarative. Defined Lifecycle.

1

Single Pane View of Addons

We want to see what cluster services are installed and it's status in one view

2

Managed Addon Lifecycle

Encapsulate Pre-install, Install and Delete workflows as part of each addon

3

Dependency Management

Define relationship dependencies between addons, order of installation, dependent status

4

Deployment Agnostic

Use your favorite deployer, helm, ship, kubectl, etc.

Addon Lifecycle

Prereqs

Create the
Addon prerequisites
e.g, cloud resources,
rbac, etc.



Install/Update

Install or update an
Addon configuration.



Validate

Validate the
components of an
Addon.

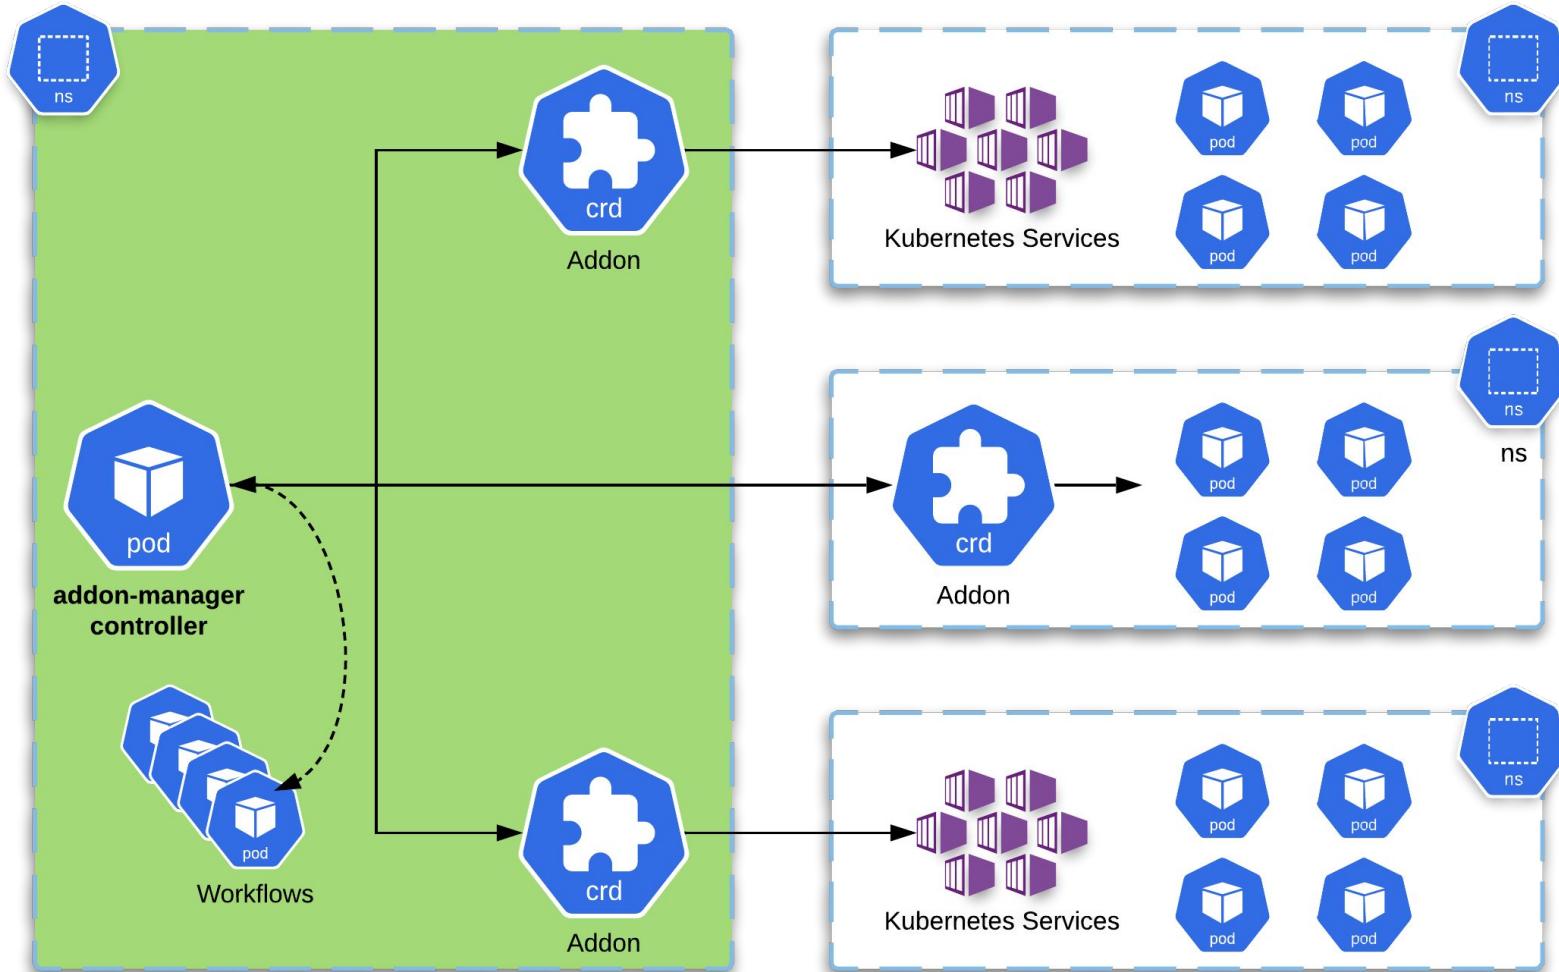


Delete

Delete the
Addon and it's
prereqs.

The background of the slide features a solid blue color with a subtle, organic white noise texture. This texture consists of fine, irregular patterns of white lines and dots that create a sense of depth and movement, resembling ripples in water or a microscopic view of a material's internal structure.

Demo



```
apiVersion: addonmgr.keikoproj.io/v1alpha1
kind: Addon
metadata:
  name: cluster-autoscaler
  namespace: addon-manager-system
spec:
  pkgChannel: cluster-autoscaler
  pkgName: cluster-autoscaler-addon
  pkgVersion: v0.1
  pkgType: composite
  pkgDescription: "cluster-autoscaler"
params:
  namespace: addon-cluster-autoscaler-ns
  context:
    clusterName: "my-example.cluster.k8s.local"
    clusterRegion: us-west-2
lifecycle:
  prereqs:
    template: |
      apiVersion: argoproj.io/v1alpha1
      kind: Workflow
      spec:
        ...
install:
  template: |
    apiVersion: argoproj.io/v1alpha1
    kind: Workflow
    spec:
      ...
delete:
  template: |
    apiVersion: argoproj.io/v1alpha1
    kind: Workflow
    spec:
      ...
```

Addon API

Descriptive package declaration

pkg* describes the unique packaging and version of an addon

params defines the injected params into Argo workflows.

lifecycle are the optional Argo Workflow templates you can define to hook into the lifecycle of an Addon

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
spec:
entrypoint: entry
serviceAccountName: addon-manager-workflow-installer-sa
templates:
- name: entry
  steps:
    - - name: prereq-resources
      template: submit
      arguments:
        artifacts:
          - name: doc
            path: /tmp/doc
            raw:
              data: |
                apiVersion: v1
                kind: Namespace
                metadata:
                  name: "{{workflow.parameters.namespace}}"
    --- 
      apiVersion: v1
      kind: ServiceAccount
      metadata:
        labels:
          k8s-addon: cluster-autoscaler.addons.k8s.io
          k8s-app: cluster-autoscaler
        name: cluster-autoscaler
        namespace: "{{workflow.parameters.namespace}}"
    --- 
      apiVersion: rbac.authorization.k8s.io/v1beta1
      kind: ClusterRole
      metadata:
        name: cluster-autoscaler
      labels:
        k8s-addon: cluster-autoscaler.addons.k8s.io
        k8s-app: cluster-autoscaler
```

Workflow API

Descriptive workflow step declaration

entrypoint describes the first step of the workflow

templates are named sequence of steps that a workflow is comprised of

steps a list of references to actions that will be executed as part of this workflow

<https://github.com/argoproj/argo>

Instance-Manager

How do I manage worker nodes?

Multitenant Clusters

The user requirements

1

Dedicated Nodes

No one likes a noisy neighbor.

2

Different Scales

Some service teams needed more nodes others needed less.

3

Different Node Types

Some service teams needed memory optimized, others needed compute, or even GPU.

Worker Node Management

The platform challenges

1

Frequent upgrades

Rotating nodes running production workloads isn't always easy, but tend to happen often.

2

Node group ownership

Kubernetes has no concept of grouping for nodes, similar to an autoscaling group, also node objects are not namespaced.

3

Multiple control planes

Changes to nodes have to be made through the cloud provider account, not directly from the cluster.

The Instance Group Lifecycle

Create

Create the instance group.



Update

Modify the instance group configuration.



Upgrade

Rotate the instances in the group.

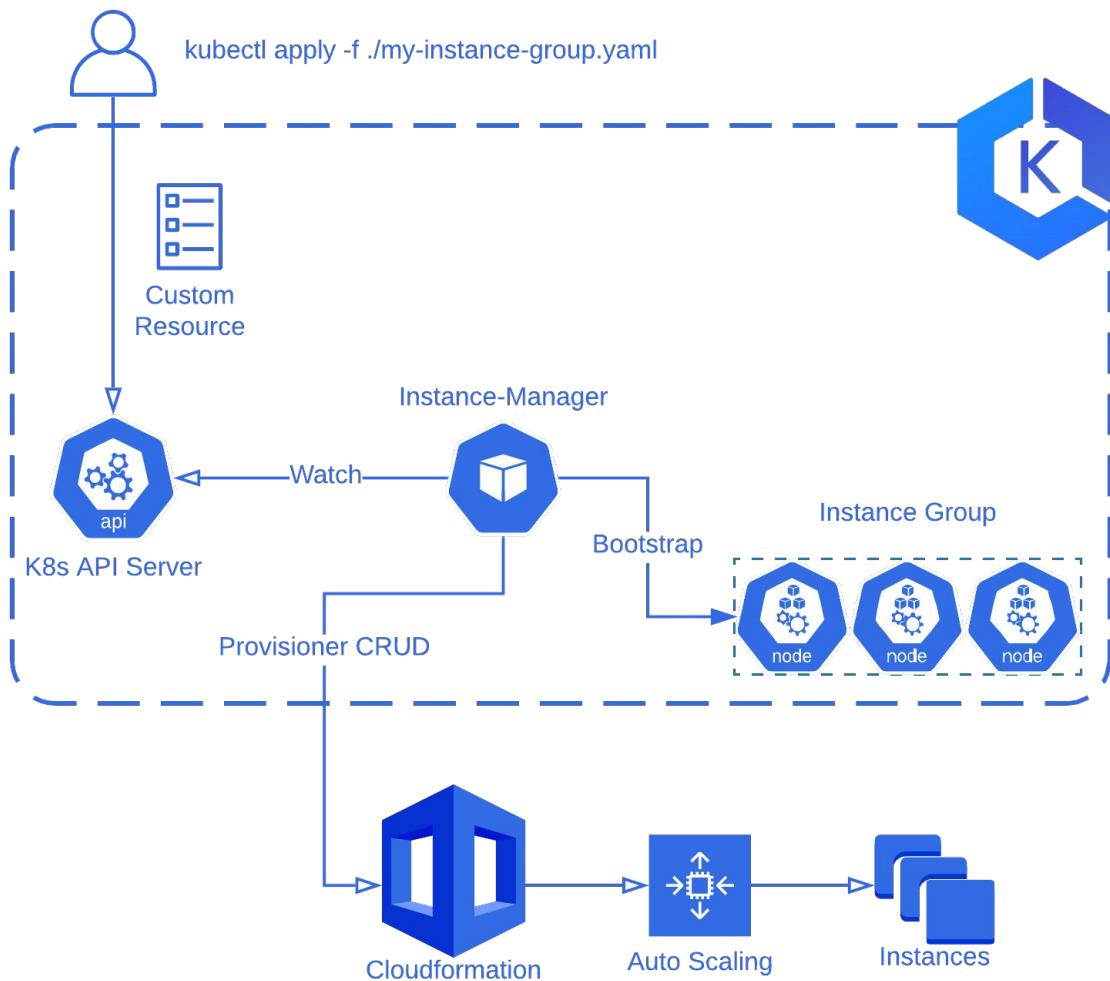


Delete

Delete the instance group.

The background of the slide features a solid blue color with a subtle, organic white noise texture. This texture consists of fine, irregular patterns of white lines and dots that create a sense of depth and movement, resembling ripples in water or a microscopic view of a material's surface.

Demo



```
apiVersion: instancemgr.keikoproj.io/v1alpha1
kind: InstanceGroup
metadata:
  name: hello-world
  namespace: instance-manager
spec:
  strategy:
    type: rollingUpdate
  rollingUpdate:
    minInstancesInService: 1
    maxBatchSize: 1
    pauseTime: PT5M
  provisioner: eks-cf
  eks-cf:
    maxSize: 6
    minSize: 3
    configuration:
      bootstrapArguments: '--node-labels=some-label=true'
      clusterName: my-eks-cluster
      subnets:
        - subnet-abb1c79de3EXAMPLE
        - subnet-cdd1c79de3EXAMPLE
        - subnet-eff1c79de3EXAMPLE
      keyPairName: my-key-pair
      image: ami-fe2dc743aEXAMPLE
      instanceType: m3.medium
      volSize: 80
      securityGroups:
        - sg-a04adb633bEXAMPLE
      tags:
        - key: my-key
          value: my-value
```

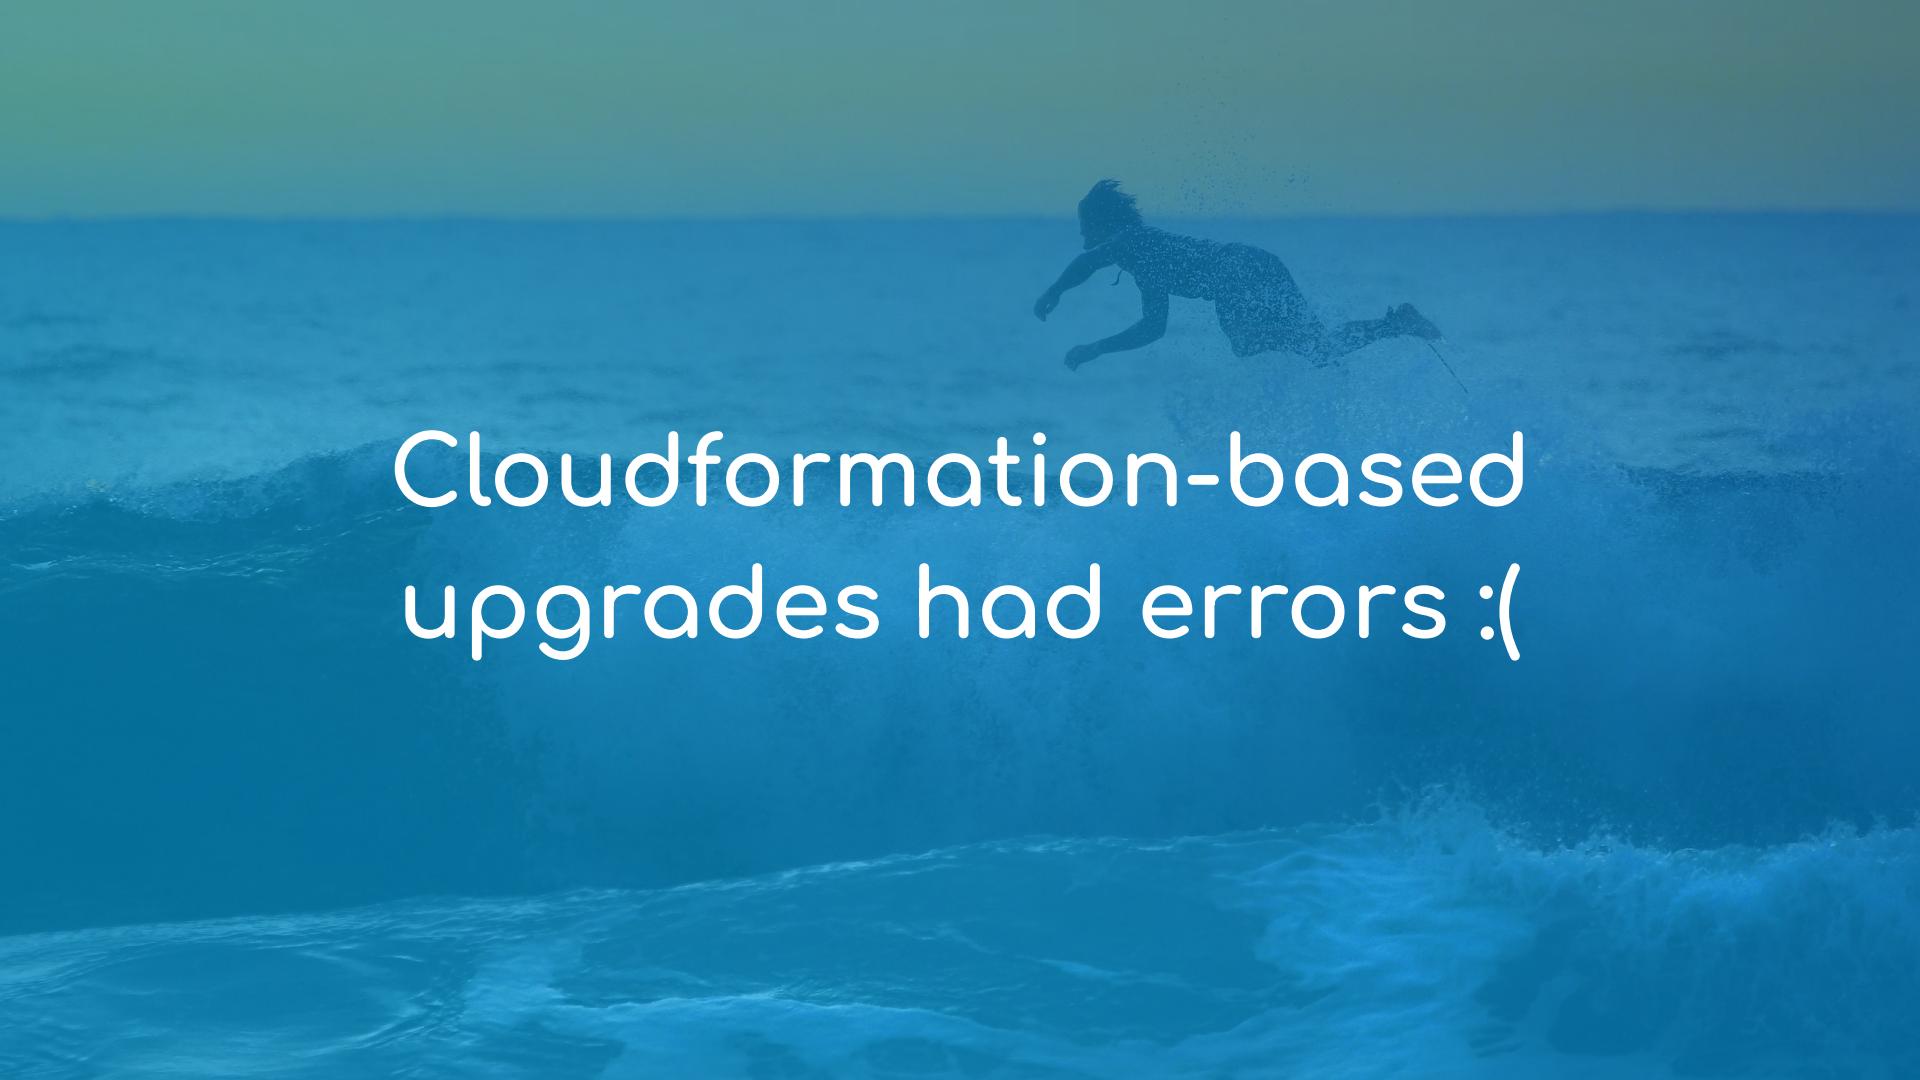
InstanceGroup API

Basic instance group with a Rolling Update strategy

strategy defines the upgrade behavior.

provisioner defines the type of instance group.

eks-cf or the named provisioner defines the configuration of the instance group.



Cloudformation-based
upgrades had errors :(

Upgrade-Manager

What does it do?

Custom Resource for Kubernetes-Aware Rolling Upgrade

Drains an instance in a scaling group

Wait for drain completion

Terminate EC2 Instance

Post-terminate wait

Move to next instance

```
apiVersion: upgrademgr.keikoproj.io/v1alpha1
kind: RollingUpgrade
metadata:
  name: my-rolling-upgrade
  namespace: instance-manager
spec:
  postDrainDelaySeconds: 30
  nodeIntervalSeconds: 300
  region: us-west-2
  asgName: my-auto-scaling-group
```

More Keiko projects!

- 1 Lifecycle-Manager
Graceful EC2 terminations using AWS lifecycle hooks
- 2 Governor
Terminates failing nodes and stuck pods
- 3 Active-Monitor
Deep monitoring of cluster capabilities
- 4 Minion-Manager
Cost savings by using spot instances
- 5 Kube-Forensics
Captures a running container state for security analysis

github.com/keikoproj